

[Start Here](#) [Blog](#) [Books](#) [About](#) [Contact](#)

Search...



How to Create an ARIMA Model for Time Series Forecasting with Python

by [Jason Brownlee](#) on [January 9, 2017](#) in [Time Series](#)



A popular and widely used statistical method for time series forecasting is the ARIMA model.

ARIMA is an acronym that stands for AutoRegressive Integrated Moving Average. It is a class of model that captures a suite of different standard temporal structures in time series data.

In this tutorial, you will discover how to develop an ARIMA model for time series data with Python.

After completing this tutorial, you will know:

- About the ARIMA model the parameters used and assumptions made by the model.
- How to fit an ARIMA model to data and use it to make forecasts.
- How to configure the ARIMA model on your time series problem.

Let's get started.

Autoregressive Integrated Moving Average Model

An [ARIMA model](#) is a class of statistical models for analyzing and forecasting time series data.

It explicitly caters to a suite of standard structures in time series data, and as such provides a simple yet powerful method for making skillful time series forecasts.

ARIMA is an acronym that stands for AutoRegressive Integrated Moving Average. It is a generalization of the simpler AutoRegressive Moving Average and adds

[Get Your Start in Machine Learning](#)

This acronym is descriptive, capturing the key aspects of the model itself. Briefly, they are:

- **AR: Autoregression.** A model that uses the dependent relationship between an observation and some number of lagged observations.
- **I: Integrated.** The use of differencing of raw observations (e.g. subtracting an observation from an observation at the previous time step) in order to make the time series stationary.
- **MA: Moving Average.** A model that uses the dependency between an observation and a residual error from a moving average model applied to lagged observations.

Each of these components are explicitly specified in the model as a parameter. A standard notation is used of ARIMA(p,d,q) where the parameters are substituted with integer values to quickly indicate the specific ARIMA model being used.

The parameters of the ARIMA model are defined as follows:

- **p:** The number of lag observations included in the model.
- **d:** The number of times that the raw observations are differenced.
- **q:** The size of the moving average window, also called the number of lagged forecast errors in the prediction equation.

A linear regression model is constructed including the independent variables that are prepared by a degree of differencing in order to make the model stationary. Structures that negatively affect the regression model.

A value of 0 can be used for a parameter, which indicates that the model is not used. The ARIMA model can be configured to perform the full AR, I, or MA model.

Adopting an ARIMA model for a time series assumes that the underlying process that generated the observations is an ARIMA process. This may seem obvious, but helps to motivate the need to confirm the assumptions of the model in the raw observations and in the residual errors of forecasts from the model.

Next, let's take a look at how we can use the ARIMA model in Python. We will start with loading a simple univariate time series.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Stop learning Time Series Forecasting the *slow way*!

Take my free 7-day email course and discover data prep, modeling and more (with sample code).

Click to sign-up and also get a free PDF Ebook version of the course.

Start Your FREE Mini-Course Now!

Get Your Start in Machine Learning

Shampoo Sales Dataset

This dataset describes the monthly number of sales of shampoo over a 3 year period.

The units are a sales count and there are 36 observations. The original dataset is credited to Makridakis, Wheelwright, and Hyndman (1998).

[Learn more about the dataset and download it from here.](#)

Download the dataset and place it in your current working directory with the filename “*shampoo-sales.csv*”.

Below is an example of loading the Shampoo Sales dataset with Pandas with a custom function to parse the date-time field. The dataset is baselined in an arbitrary date.

```
from pandas import read_csv
from pandas import datetime
from matplotlib import pyplot

def parser(x):
    return datetime.strptime('190'+x, '%Y-%m')

series = read_csv('shampoo-sales.csv', header=0,
                  parse_dates=[0], dayfirst=True,
                  squeeze=True, index_col=0)
print(series.head())
series.plot()
pyplot.show()
```

Running the example prints the first 5 rows of the data:

```
1 Month
2 1901-01-01 266.0
3 1901-02-01 145.9
4 1901-03-01 183.1
5 1901-04-01 119.3
6 1901-05-01 180.3
7 Name: Sales, dtype: float64
```

The data is also plotted as a time series with the month along the x-axis and sales figures on the y-axis.

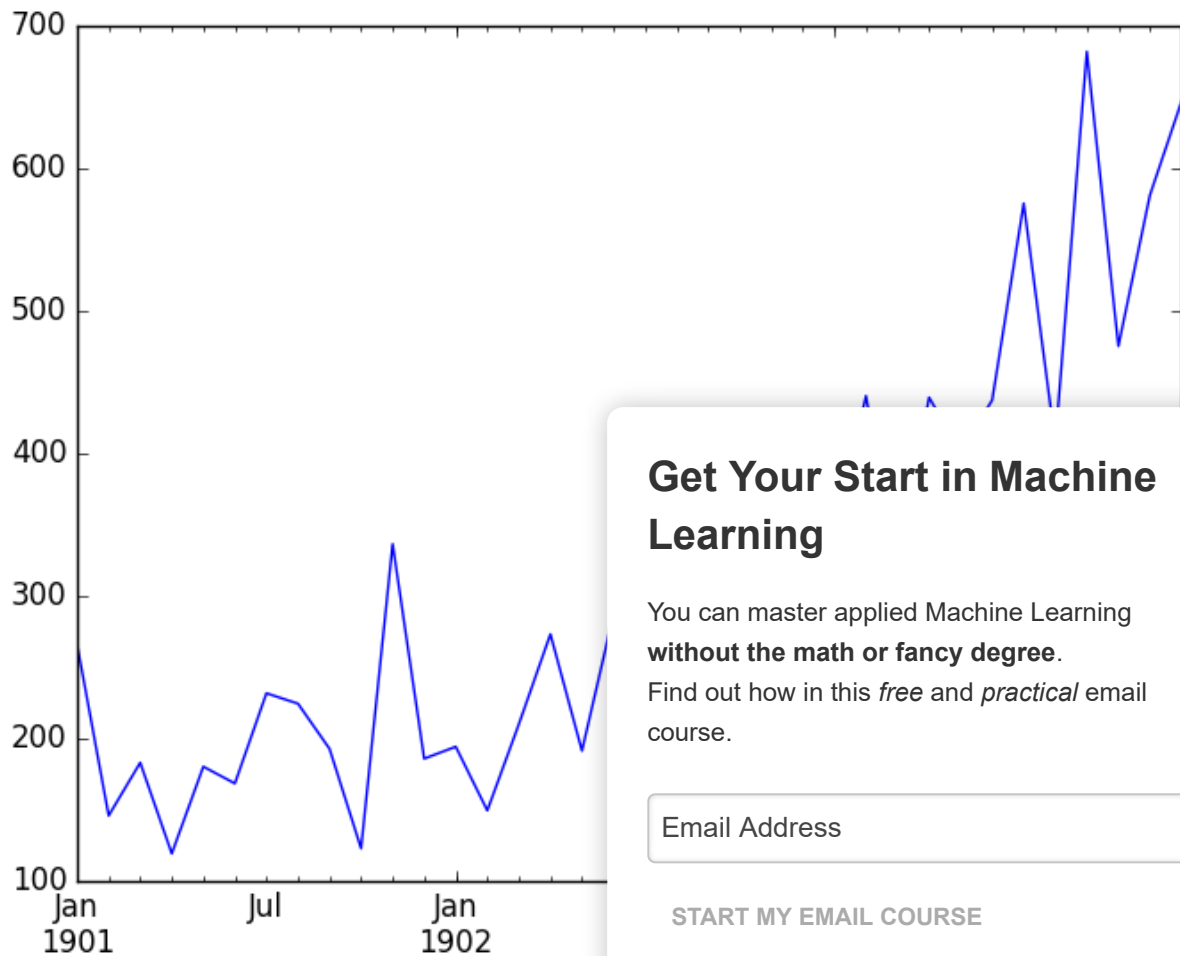
Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Shampoo Sales Dataset Plot

We can see that the Shampoo Sales dataset has a clear trend.

This suggests that the time series is not stationary and will require differencing to make it stationary, at least a difference order of 1.

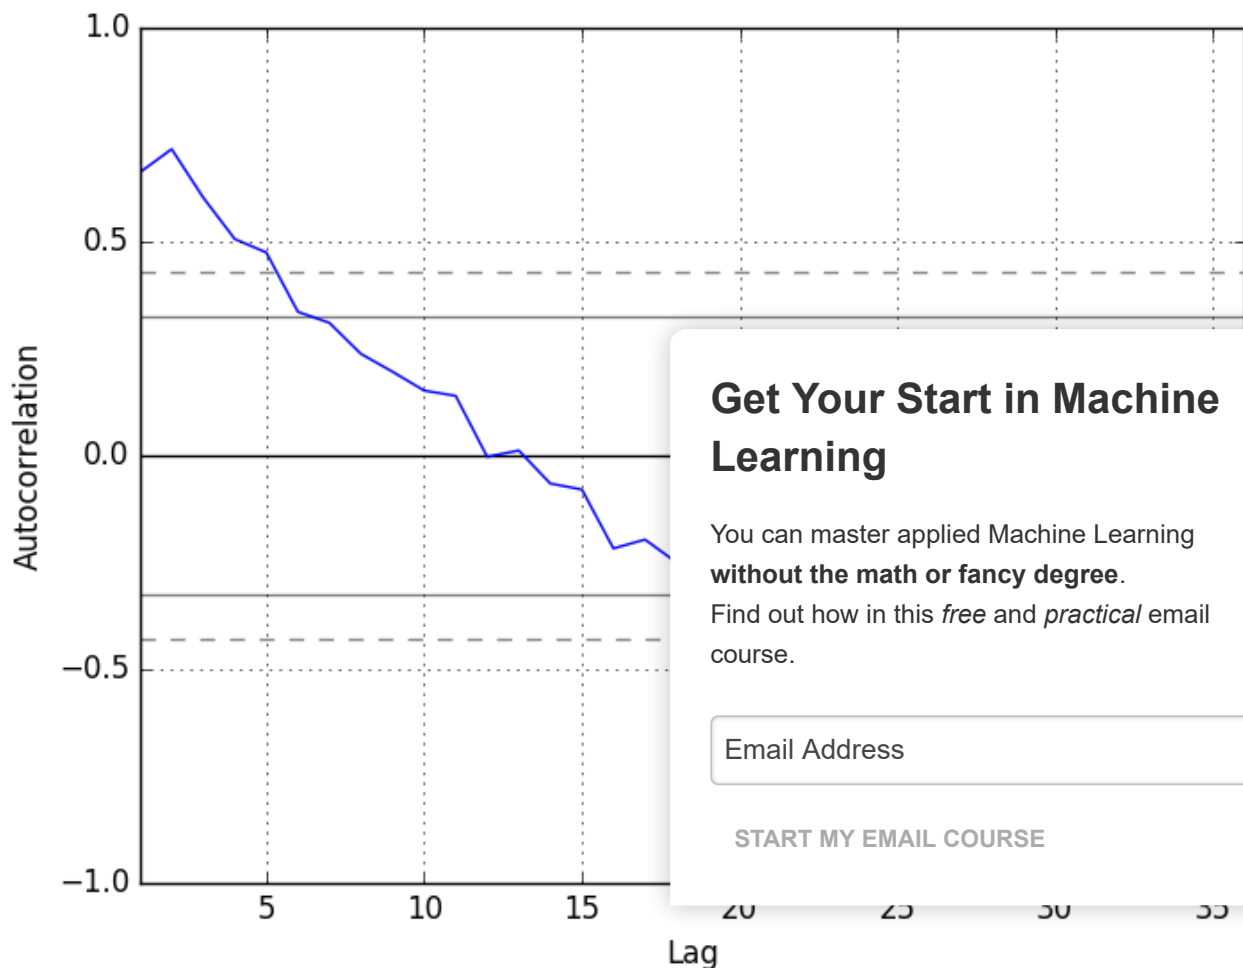
Let's also take a quick look at an autocorrelation plot of the time series. This is also built-in to Pandas. The example below plots the autocorrelation for a large number of lags in the time series.

```
1 from pandas import read_csv
2 from pandas import datetime
3 from matplotlib import pyplot
4 from pandas.tools.plotting import autocorrelation_plot
5
6 def parser(x):
7     return datetime.strptime('190'+x, '%Y-%m')
8
9 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, date_parser=parser)
10 autocorrelation_plot(series)
11 pyplot.show()
```

Running the example, we can see that there is a positive correlation with the first 10-to-12 lags that is perhaps significant for the first 5 lags.

Get Your Start in Machine Learning

A good starting point for the AR parameter of the model may be 5.



Autocorrelation Plot of Shampoo Sales Data

ARIMA with Python

The statsmodels library provides the capability to fit an ARIMA model.

An ARIMA model can be created using the statsmodels library as follows:

1. Define the model by calling `ARIMA()` and passing in the p , d , and q parameters.
2. The model is prepared on the training data by calling the `fit()` function.
3. Predictions can be made by calling the `predict()` function and specifying the index of the time or times to be predicted.

Let's start off with something simple. We will fit an ARIMA model to the entire Shampoo Sales dataset and review the residual errors.

First, we fit an `ARIMA(5,1,0)` model. This sets the lag value to 5 for autoregression, uses a difference order of 1 to make the time series stationary, and uses a model

Get Your Start in Machine Learning

When fitting the model, a lot of debug information is provided about the fit of the linear regression model. We can turn this off by setting the *disp* argument to 0.

```

1 from pandas import read_csv
2 from pandas import datetime
3 from pandas import DataFrame
4 from statsmodels.tsa.arima_model import ARIMA
5 from matplotlib import pyplot
6
7 def parser(x):
8     return datetime.strptime('190'+x, '%Y-%m')
9
10 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, do
11 # fit model
12 model = ARIMA(series, order=(5,1,0))
13 model_fit = model.fit(dis=0)
14 print(model_fit.summary())
15 # plot residual errors
16 residuals = DataFrame(model_fit.resid)
17 residuals.plot()
18 pyplot.show()
19 residuals.plot(kind='kde')
20 pyplot.show()
21 print(residuals.describe())

```

Running the example prints a summary of the fit model as well as the skill of the fit on the on the in-sample observations.

```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28

```

ARIMA Model Results						
Dep. Variable:	D.Sales	No. Obs:	190			
Model:	ARIMA(5, 1, 0)	Log Likelihood:	-417.227			
Method:	css-mle	S.D. Error:	0.183			
Date:	Mon, 12 Dec 2016	AIC:	417.227			
Time:	11:09:13	BIC:	417.227			
Sample:	02-01-1901	HQIC:	410.098			
	- 12-01-1903					
	coef	std err	z	P> z	[95.0% Conf. Int.]	
const	12.0649	3.652	3.304	0.003	4.908 19.222	
ar.L1.D.Sales	-1.1082	0.183	-6.063	0.000	-1.466 -0.750	
ar.L2.D.Sales	-0.6203	0.282	-2.203	0.036	-1.172 -0.068	
ar.L3.D.Sales	-0.3606	0.295	-1.222	0.231	-0.939 0.218	
ar.L4.D.Sales	-0.1252	0.280	-0.447	0.658	-0.674 0.424	
ar.L5.D.Sales	0.1289	0.191	0.673	0.506	-0.246 0.504	
Roots						
	Real	Imaginary	Modulus	Frequency		
AR.1	-1.0617	-0.5064j	1.1763	-0.4292		
AR.2	-1.0617	+0.5064j	1.1763	0.4292		
AR.3	0.0816	-1.3804j	1.3828	-0.2406		
AR.4	0.0816	+1.3804j	1.3828	0.2406		
AR.5	2.9315	-0.0000j	2.9315	-0.0000		

First, we get a line plot of the residual errors, suggesting that there may still be some trend information not captured by the model.


Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**


Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



Next, we get a density plot of the residual error values, suggesting the errors are Gaussian, but may not be centered on zero.



The distribution of the residual errors is displayed. The results show that indeed there is a bias in the prediction (a non-zero mean in the residuals).

1	count	35.000000
2	mean	-5.495213
3	std	68.132882
4	min	-133.296597
5	25%	-42.477935
6	50%	-7.186584
7	75%	24.748357
8	max	133.237980

Note, that although above we used the entire dataset for time series analysis, ideally we would perform this analysis on just the training dataset when developing a predictive model.

Next, let's look at how we can use the ARIMA model to make forecasts.

Rolling Forecast ARIMA Model

The ARIMA model can be used to forecast future time steps.

Get Your Start in Machine Learning

We can use the `predict()` function on the `ARIMAResults` object to make predictions. It accepts the index of the time steps to make predictions as arguments. These indexes are relative to the start of the training dataset used to make predictions.

If we used 100 observations in the training dataset to fit the model, then the index of the next time step for making a prediction would be specified to the prediction function as `start=101, end=101`. This would return an array with one element containing the prediction.

We also would prefer the forecasted values to be in the original scale, in case we performed any differencing ($d>0$ when configuring the model). This can be specified by setting the `typ` argument to the value `'levels'`: `typ='levels'`.

Alternately, we can avoid all of these specifications by using the `step_ahead` argument to make a one-step forecast using the model.

We can split the training dataset into train and test sets to evaluate the model's prediction for each element on the test set.

A rolling forecast is required given the dependence on the AR model. A crude way to perform this rolling forecast is to update the model each time a new observation is received.

We manually keep track of all observations in a list called `history` to which new observations are appended each iteration.

Putting this all together, below is an example of a rolling forecast.

```
1 from pandas import read_csv
2 from pandas import datetime
3 from matplotlib import pyplot
4 from statsmodels.tsa.arima_model import ARIMA
5 from sklearn.metrics import mean_squared_error
6
7 def parser(x):
8     return datetime.strptime('190'+x, '%Y-%m')
9
10 series = read_csv('shampoo-sales.csv', header=0, parse_dates=[0], index_col=0, squeeze=True, dtype=object)
11 X = series.values
12 size = int(len(X) * 0.66)
13 train, test = X[0:size], X[size:len(X)]
14 history = [x for x in train]
15 predictions = list()
16 for t in range(len(test)):
17     model = ARIMA(history, order=(5,1,0))
18     model_fit = model.fit(dis=0)
19     output = model_fit.forecast()
20     yhat = output[0]
21     predictions.append(yhat)
22     obs = test[t]
23     history.append(obs)
24     print('predicted=%f, expected=%f' % (yhat, obs))
25 error = mean_squared_error(test, predictions)
26 print('Test MSE: %.3f' % error)
27 # plot
```

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning

```
28 pyplot.plot(test)
29 pyplot.plot(predictions, color='red')
30 pyplot.show()
```

Running the example prints the prediction and expected value each iteration.

We can also calculate a final mean squared error score (MSE) for the predictions, providing a point of comparison for other ARIMA configurations.

```
1 predicted=349.117688, expected=342.300000
2 predicted=306.512968, expected=339.700000
3 predicted=387.376422, expected=440.400000
4 predicted=348.154111, expected=315.900000
5 predicted=386.308808, expected=439.300000
6 predicted=356.081996, expected=401.300000
7 predicted=446.379501, expected=437.400000
8 predicted=394.737286, expected=575.500000
9 predicted=434.915566, expected=407.600000
10 predicted=507.923407, expected=682.000000
11 predicted=435.483082, expected=475.300000
12 predicted=652.743772, expected=581.300000
13 predicted=546.343485, expected=646.900000
14 Test MSE: 6958.325
```

A line plot is created showing the expected values (blue) and predicted values (red). We can see the values show some trend and are in the same range (around 300-700).

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Get Your Start in Machine Learning



ARIMA Rolling Forecast Line Plot

The model could use further tuning of the p , d , and maybe even the q parameters.

Configuring an ARIMA Model

The classical approach for fitting an ARIMA model is to follow the [Box-Jenkins Methodology](#).

This is a process that uses time series analysis and diagnostics to discover good parameters for the ARIMA model.

In summary, the steps of this process are as follows:

1. **Model Identification.** Use plots and summary statistics to identify trends, seasonality, and autoregression elements to get an idea of the amount of differencing and the size of the lag that will be required.
2. **Parameter Estimation.** Use a fitting procedure to find the coefficients of the regression model.
3. **Model Checking.** Use plots and statistical tests of the residual errors to determine the amount and type of temporal structure not captured by the model.

Get Your Start in Machine Learning

The process is repeated until either a desirable level of fit is achieved on the in-sample or out-of-sample observations (e.g. training or test datasets).

The process was described in the classic 1970 textbook on the topic titled [Time Series Analysis: Forecasting and Control](#) by George Box and Gwilym Jenkins. An updated 5th edition is now available if you are interested in going deeper into this type of model and methodology.

Given that the model can be fit efficiently on modest-sized time series datasets, grid searching parameters of the model can be a valuable approach.

Summary

In this tutorial, you discovered how to develop an ARIMA model.

Specifically, you learned:

- About the ARIMA model, how it can be configured
- How to perform a quick time series analysis using
- How to use an ARIMA model to forecast out of sample


Do you have any questions about ARIMA, or about this tutorial? Ask your questions in the comments below and I will do my best to answer them.

Get Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree**. Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Want to Develop Time Series Forecasts with Python?



Develop Your Own Forecasts in Minutes

...with just a few lines of python code

Discover how in my new Ebook:

[Introduction to Time Series Forecasting With Python](#)

It covers **self-study tutorials** and **end-to-end projects** on topics like: *Loading data, visualization, modeling, algorithm tuning*, and much more...

Finally Bring Time Series Forecasting to Your Own Projects

Skip the Academics. Just Results.

[Click to learn more.](#)

Get Your Start in Machine Learning